

# Advanced Machine Learning - Classification

## Midterm Project

November 2023

### 1 Instructions

- The entire project will be graded and will contribute 30% of the final grade.
- It is required to submit notebooks only, you should not write a final report. Nevertheless, the notebooks you submit should contain comments and explanations of your procedures.
- I will not grade you on the final quality of the classification tasks, but I will grade i) correctness of the Python implementation, ii) quality of the dataset visualization, iii) quality of the comments and findings, iv) scientific correctness of the performed steps.
- Please submit only one zip file containing all the scripts you wrote.
- Submit the zip file on iCorsi. The deadline is **November 19, 23:59**.
- You can talk about the task to your fellow students, but please do not share your implementation/code. Plagiarism in the source code or the report will be penalized with a grade of 0 and might be reported to the faculty.
- You are absolutely welcome to ask me anything about things that are unclear to you in this assignment. Feel free to contact me for feedback or suggestions.

## 2 Tasks

- If not specified, for evaluation of your results use the metric you think is more appropriate among those we saw in the lectures.

### 2.1 Task 1

In the following Tasks you are going to work with the digits dataset in sklearn and the Fashion MNIST dataset from `tensorflow.keras`. Use the tool you prefer (e.g., LDA, t-SNE) to perform a preliminary visualization of the two datasets.

### 2.2 Task 2

Run a multilayer perceptron (a densely connected feed forward neural network) with two hidden layers and rectified linear nonlinearities (i.e., RELU) on the digits dataset from sklearn using the `tensorflow.keras Sequential` interface. Include code for selecting  $l_2$ -norm regularization strength and number of hidden units using `GridSearchCV` and evaluation on an independent test-set.

#### Bonus points

*Transfer learning* - Reuse an existing architecture and pre-trained weights from keras (<https://keras.io/api/applications/>). Compare retraining only the densely connected layers with fine-tuning the whole network.

### 2.3 Task 3

Train a multilayer perceptron (fully connected) on the Fashion MNIST dataset using the traditional train/test split as given by `fashion_mnist.load_data` in `tensorflow.keras`. Use a separate 10000 samples (from the training set) for model selection and to compute learning curves (accuracy vs epochs, not accuracy vs n.samples). Compare a “vanilla” model with a model using drop-out and evaluate if using drop-out allows you to learn a bigger network. Then, compare to a model using batch normalization. Visualize learning curves for all models.

#### Bonus points

Augment the data using rotations, mirroring and possibly other transformations. How much can you improve your original model by data augmentation?

## 2.4 Task 4

Perform calibration on one of the previously implemented models of your choice. Plot the Reliability Diagram before and after the calibration procedure.

## 3 Suggestions

- You can use `StratifiedShuffleSplit` to create a single train-validation split for use with `GridSearchCV`. Use of `GridSearchCV` might not be the best option for any task but Task 2, though.
- Preprocess the images before visualization and/or training a model.
- Test your code on a small part of the data before training the model. You don't want your code to fail on a print statement after waiting for the network to train.